

Chen Borun

+ (86) 15217720601 ◇ chenborun1101@gmail.com ◇ github.com/PorunC

PROFESSIONAL SUMMARY

3 years of backend development experience, focused on bringing Large Language Models into production-grade business tools. Led three 0-to-1 projects, including the AI Weekly Report System and the Code Intelligence Platform, with proven ability to independently deliver end-to-end system design, full-stack development and continuous delivery.

WORK EXPERIENCE

CMB Network Technology Co., Ltd.

2023.7 – Present

Backend Software Engineer / Core Application Development Team

- **AI Weekly Report System (Independently Owned)**

[Situation] The bank's operations weekly report relied on manual data collection and aggregation across multiple source systems, taking 8–10 hours per cycle, which was inefficient and error-prone.

[Task] Build a fully automated, LLM-driven weekly report system.

[Action] Designed and built from scratch an end-to-end pipeline covering multi-source heterogeneous data ingestion, LLM-based segmented analysis, report assembly and delivery, with built-in retry, fault tolerance and transactional guarantees.

[Result] Reduced weekly report turnaround from 8–10 hours to 20 minutes, saving roughly 50 person-days per year, and the system has been running stably in production since launch.

- **CodeWiki – Code Intelligence Platform (Independently Owned)**

[Situation] Cross-module code understanding required engineers to read source files line by line; new joiners typically needed weeks to ramp up, limiting overall R&D efficiency.

[Task] Build a code-understanding and knowledge-management platform to lower the team's cognitive cost when working with the codebase.

[Action] Built a multi-language AST parsing engine and a multi-dimensional dependency graph from scratch; implemented semantic-level code retrieval based on GraphRAG, driving the LLM to automatically generate source-level documentation; delivered the full stack independently (Python/FastAPI + TypeScript/React) with multiple interaction modes.

[Result] 50+ daily queries on average; helped the team locate cross-module dependencies and implementation details much faster, and was adopted in day-to-day code reviews and new-hire onboarding.

- **AI Agent Long-Term Memory & Context Compression System (Independently Owned)**

[Situation] AI agents lacked stable long-term memory in long-running tasks and multi-turn conversations; historical context and tool-call logs kept expanding, causing reasoning accuracy to drop and token cost to climb.

[Task] Design and ship a production-grade agent memory system that improves context utilization, recall accuracy and execution stability for long-horizon tasks.

[Action] Designed an L0–L3 layered memory architecture from scratch, managing raw conversations, structured summaries, semantic memory and long-term knowledge in tiers; built a memory retrieval pipeline based on vector search, BM25 full-text search and Hybrid RRF fusion ranking, with dual back-ends for local SQLite and cloud vector databases; designed a context-offloading and tool-log compression mechanism that

summarizes, traces and recovers long-running task execution; complemented it with automatic scheduling, session cleanup, metrics reporting and millisecond-level performance timing to keep the system observable and maintainable.

[Result] The system has been deployed and adopted by multiple business channels: token consumption reduced by 61%, task pass rate improved by 52% (relative), and long-term memory accuracy lifted from 48% to 76%.

EDUCATION

Shenzhen University

2019.9 – 2023.6

B.Eng. in Software Engineering

OTHERS

- English: CET-6, able to read and write English technical documentation
- Blog: misaka-9982.com